


# Deep Learning on Graphs

## Part 2

*Michael Kenning, Stavros Georgousis*

# Spectral Convolution cont.

- Same objective as spatial convolution
  - (Theoretically) no problem with neighbourhood selection.
  - Signal diffused from central node.
- Cons -> Computationally expensive



Source: Xu et al. 2019

# Spectral Convolution cont.

Idea of spectral convolution:

1. Project graph signal to spectral domain
2. Filter in spectral domain.
3. Get a new feature representation in vertex (spatial) domain.

$$\tilde{f} = \Phi^{\top} f$$

$$f = \Phi \tilde{f}$$

## Example: Spectral Convolution (Bruna et. al., 2014)

- Fourier transform  $\rightarrow$   $\mathbf{U}$  = eigenvectors of  $\mathbf{L}$ :  $\hat{\mathbf{L}} = \mathbf{U}\hat{\mathbf{\Lambda}}\mathbf{U}^\top$

- Convolution defined as:

$$f \star g = \mathbf{U} \left( (\mathbf{U}^\top f) \odot (\mathbf{U}^\top g) \right)$$

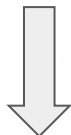
- Layer implementation:  $f_{l+1,j} = h \left( \mathbf{U} \sum_{i=0}^{c-1} \left( \Theta_{l,i,j} \mathbf{U}^\top f_{l,i} \right) \right)$

- Computationally expensive:  $O(n^2)$

# Example: ChebNet (Defferard et al., 2016)

- Calculation of eigenvalue decomposition computationally expensive
- Solution -> approximate the Fourier transformed signal (Chebyshev polynomials)

Approximate !


$$\mathbf{g}_\theta(\mathbf{\Lambda}) \approx \sum_{i=0}^{k-1} \theta_i T_i(\bar{\mathbf{\Lambda}}),$$
$$\mathbf{g}_\theta(\mathbf{L}) \approx \sum_{i=0}^{k-1} \theta_i T_i(\bar{\mathbf{L}}) f$$

$$f_{l+1,j} = h \left( \sum_{i=0}^{c-1} \mathbf{g}_\theta(\bar{\mathbf{L}}) f_{l,i} \right)$$

# Example: Graph Convolution Network (GCN) (Kipf et al., 2017)

Where we left off:

Approximate !

$$\mathbf{g}_\theta(\mathbf{L}) \approx \sum_{i=0}^{k-1} \theta_i T_i(\bar{\mathbf{L}}) f$$

$$f \star \mathbf{g} \approx \theta'_0 f + \theta'_1 (\mathbf{L} - \mathbf{I}_n) f = \theta'_0 f - \theta'_1 \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} f.$$

$$f \star \mathbf{g} \approx \theta \left( \mathbf{I}_n + \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} \right) f = \theta \mathbf{L} f$$

Renormalisation trick

$$f_{l+1} = h \left( \bar{\mathbf{D}}^{-1/2} \bar{\mathbf{A}} \bar{\mathbf{D}}^{-1/2} f_l \Theta_l \right)$$

$$\mathbf{I}_n + \bar{\mathbf{D}}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} \rightarrow \bar{\mathbf{D}}^{-1/2} \bar{\mathbf{A}} \bar{\mathbf{D}}^{-1/2}$$

$$\bar{\mathbf{A}} = \mathbf{A} + \mathbf{I}_n,$$

- Computationally efficient (well... relatively)
- Neighbourhood order: 1
- Basically the baseline standard
- Curious case (could be said it isn't spectral)

# Graph Pooling

Global pooling (readout):


- Summarises features for graph level tasks
- No equivalent in CNNs

Hierarchical pooling:

- Analogous to standard pooling in CNNs
- Dilates receptive field
- Can improve performance using hierarchical representation

# Global pooling (readout layer)


- Min, max, mean, sum pooling
- Combinations of above mean + max pool (Ying et al., 2018)
- Interesting example: SortPool (Zhang et al., 2019)






# Hierarchical Pooling

- Subsampling based
  - Pros -> Efficient
  - Cons -> Discard information, potentially isolated vertices
- Clustering based
  - Common caveat: pooled graph dense



# Subsampling based pooling

Examples: Top-K pooling (Graph U-Nets, SAGPool)




# Clustering based pooling

Common approach learning a clustering assignment matrix,  $S$

$$\mathbf{A}_{\text{pool}} = \mathbf{S}^T \mathbf{A} \mathbf{S}$$
$$\mathbf{A}_{\text{unpool}} = \mathbf{S} \mathbf{A}_{\text{pool}} \mathbf{S}^T$$

$S \rightarrow$  truncated matrix  
of weights from  
original nodes to supernodes



Source: Ying et al. 2019

# Example: Edge Pooling

Selects edges to be contracted based on edge scores.

- Pros: Retains sparsity

